

FountainAI Platform Overview

What is FountainAI?

FountainAI is a comprehensive platform that combines large language models (LLMs) with a suite of specialized services to enable advanced AI reasoning, planning, and knowledge management. Its mission is to provide an AI orchestration environment where an LLM can **plan tasks, call external functions (tools), reflect on results, and analyze changes in knowledge** over time. The platform manages "semantic artifacts" – such as knowledge baselines, drift documents, reflections, and function definitions – organized into **corpora** (contexts or workspaces) ¹. By integrating function-calling LLMs with persistent semantic memory, FountainAI supports use cases like dynamic tool use, step-by-step task planning, automated self-reflection, and drift analysis to track how information evolves.

Tools Factory

The **Tools Factory Service** is responsible for registering and managing external tools (functions) that the AI can invoke during its reasoning process ². In FountainAI, a "tool" is essentially a function defined by an OpenAPI specification. The Tools Factory takes in these definitions and stores them in a shared index (backed by a Typesense search engine) so that they can be discovered and called by other services (notably the Function Caller) ³. Key capabilities of the Tools Factory include:

- **Registering New Tools:** Clients can submit an OpenAPI document (containing one or multiple function definitions) via the `/tools/register` endpoint. The service extracts each operation (function) and saves its metadata (name, description, parameters schema, HTTP endpoint to call, etc.) for later use ⁴ ⁵. Each tool is associated with a unique `operationId` (function identifier) and can be tagged to a specific corpus context (via a `corpusId` parameter) for namespacing ⁶.
- **Listing Available Tools:** The `/tools` endpoint allows retrieval of all registered functions, with pagination support ⁷ ⁸. This lets the system or users query what tools are currently available for the AI to use.

By centralizing function definitions, the Tools Factory enables FountainAI to dynamically expand its capabilities. When new APIs or functions are registered, the AI can discover and invoke them as needed to fulfill user objectives.

Function Caller

The **Function Caller Service** is the runtime component that **executes** the tools/functions registered in the platform ⁹. It serves as a bridge between the LLM's function-calling intentions and actual HTTP API calls or other function executions. In practice, when the LLM (via the Planner) decides to use a function, the Function Caller handles invoking that function and returning the results. Key functionalities of this service include:

- **Function Invocation:** The core feature is an **invoke endpoint** (`POST /functions/{function_id}/invoke`) that takes a function identifier and input parameters, then performs the corresponding HTTP request or operation and returns the result ¹⁰ ¹¹. This allows the LLM

to trigger any available tool by name (the `function_id` corresponds to the `operationId` defined when the tool was registered).

- **Function Management:** The Function Caller can **list all registered functions** (`GET /functions`) and provide **details of a specific function** (`GET /functions/{function_id}`)^{12 13}. This mirrors the data from Tools Factory, ensuring the AI agent knows what functions exist and how to call them. Each function's metadata includes its human-readable name, description, HTTP method and path, and a schema for expected parameters^{14 15}.

By acting as a “dynamic *operationId-to-HTTP call mapping factory*”, the Function Caller enables **LLM-driven function orchestration**¹⁶. In other words, the AI can plan a series of function calls (by their IDs) and rely on this service to carry out the actual operations and retrieve results, which can then be fed back into the LLM's reasoning process.

Bootstrap Service

The **Bootstrap Service** handles the **initialization and setup of a new corpus** (a working context for an AI agent) and streamlines the process of seeding it with default information¹⁷. Think of it as a one-stop orchestrator for starting an AI session with all necessary defaults (roles, baseline knowledge, etc.) and for adding new content with immediate analysis. Its key responsibilities and endpoints include:

- **Corpus Initialization:** The endpoint `POST /bootstrap/corpus/init` creates a new empty corpus and immediately seeds it with some defaults¹⁸. Under the hood, this call performs multiple steps: (1) it calls the Baseline Awareness service to create a new corpus entry (essentially reserving an ID for a new knowledge base); (2) it seeds the corpus with five default GPT **roles** via another endpoint; and (3) it enqueues a default **reflection** (called “role-health-check”) into the corpus^{19 20}. The default roles are pre-defined prompts or personas for the GPT model, ensuring the agent starts with some baseline behaviors or perspectives.
- **Seeding GPT Roles:** The `POST /bootstrap/roles/seed` endpoint explicitly populates an existing corpus with the five standard GPT role prompts^{21 22}. These roles might represent different viewpoints or skills (for example, Analyst, Critic, Explainer, etc., although the specific roles are defined in the system) that the AI can adopt. Seeding roles ensures the AI has a multi-faceted starting point for reasoning.
- **Default Reflection and Role Promotion:** The `POST /bootstrap/corpus/reflect` can enqueue a default reflection (the “role-health-check”) at any time for a given corpus²³, prompting the system to generate a reflective analysis (often a Q&A or diagnostic about the current roles’ performance). Additionally, `POST /bootstrap/roles/promote` will take the latest reflection (specifically the result of a role health check) and **convert it into a new GPT role**^{24 25}. This means if the system's reflection identifies a useful new perspective or rule, that insight can be turned into a permanent role prompt for future interactions.
- **Adding Baselines with Analysis:** One of the most powerful features is `POST /bootstrap/baseline`, which allows the user to add a new **baseline document** to the corpus along with triggering immediate analysis²⁶. A *baseline* could be a chunk of knowledge or data (text or JSON) that the user considers the ground truth or starting content. When this endpoint is called, the service will:
 - Store the baseline content in the Awareness/Persistence layer²⁷.
 - Kick off **real-time analysis streams** for *drift* and *patterns*. Specifically, it opens two server-sent event (SSE) channels – one that streams a **drift analysis** (comparing the new baseline with prior state to detect what has changed or “drifted”) and another that streams a **narrative patterns analysis** (identifying themes or patterns in the content)²⁸. These analyses are persisted as special entries (tagged with the baseline ID plus “-drift” and “-patterns”) once completed.

- Schedule follow-up background jobs to update the **history aggregation** and **semantic arc** for the corpus ²⁹. The history is an accumulation of all semantic changes, and the semantic arc is a synthesized overview or storyline of the corpus's knowledge over time.

All these steps happen in one call to `/bootstrap/baseline`, providing a convenient way to enrich the corpus and immediately understand new data ³⁰. The Bootstrap Service thus ties together the other services (Awareness, Roles, Persistence) to simplify starting and updating an AI agent's knowledge base.

Baseline Awareness Service

The **Baseline Awareness Service** manages the **knowledge state and semantic analytics** for each corpus ³¹. It is essentially the brain of FountainAI for tracking what the AI knows (baselines), how it changes (drift), what patterns emerge, and what reflections have been made. Its core functions and endpoints include:

- **Corpus Management:** The service can initialize a new **corpus** (`POST /corpus/init`), which creates a fresh container for semantic content identified by a corpus ID ³². All subsequent data (baselines, drifts, etc.) will be tied to this corpus.
- **Baselines and Drift:** You can add a **baseline** text or document to the corpus via `POST /corpus/baseline` ³³. A baseline might be an initial snapshot of information or a significant update. Complementary to baselines, `POST /corpus/drift` allows adding a **drift document** ³⁴. A drift document typically represents changes or differences from a previous baseline – essentially capturing how the state has shifted. The Awareness service likely computes or stores these drifts (in practice, as noted, the Bootstrap service triggers drift analysis automatically when adding a baseline).
- **Narrative Patterns:** Through `POST /corpus/patterns`, the service can store **narrative patterns** related to the corpus ³⁵. Narrative patterns analysis might extract common themes, story arcs, or other structural patterns from the content, which help in understanding the overall narrative or evolution of the data.
- **Reflections:** The Awareness service records **reflections** – which are question/answer pairs or insights generated by the AI about the corpus. Clients can add a reflection via `POST /corpus/reflections` ³⁶, providing a prompt/question and the AI's reflective answer. The service stores these as part of the corpus's knowledge. You can also **list all reflections** in a corpus with `GET /corpus/reflections/{corpus_id}` ³⁷ ³⁸. Reflections are a form of automated semantic reasoning where the AI thinks about what it knows or about its own processes, aiding in planning and self-correction.
- **History and Semantic Arc:** The service keeps a **history** of all changes in the corpus. `GET /corpus/history/{corpus_id}` returns the change log (baselines added, drifts, reflections made, etc.) over time ³⁹ ⁴⁰. Moreover, it can generate a **summarized history** or "semantic arc" for the corpus: `GET /corpus/summary/{corpus_id}` gives a semantic summary of the history ⁴¹ ⁴², and `GET /corpus/semantic-arc` produces a more analytical construct of how the narrative/knowledge evolved ⁴³ ⁴⁴. The *semantic arc* is essentially an insight into the corpus's journey – highlighting key developments, shifts in theme or sentiment, and overall progress of understanding.

In summary, the Baseline Awareness Service is crucial for **semantic reasoning and memory**. It not only stores raw data (baselines and updates) but also layers on interpretation – tracking differences (drift), extracting patterns, and allowing the system to reflect and summarize. This ensures the FountainAI agent remains *aware of its knowledge state* and can reason about its own knowledge (meta-cognition).

Planner Service

The **Planner Service** is the component that enables **LLM-driven planning and orchestration** of tasks ⁴⁵. When a user provides a high-level goal or problem, the Planner service is responsible for breaking it down (with the help of the LLM) into a sequence of actionable steps, usually in the form of function calls (tools to use) and then executing those steps. The Planner's key features include:

- **Generating Plans (Reasoning):** The endpoint `POST /planner/reason` takes a user's **objective** (a description of what the user wants to achieve) and produces a **step-by-step plan** to accomplish it ⁴⁶ ⁴⁷. Under the hood, the Planner likely engages an LLM (via the LLM Gateway) along with context about available tools and relevant knowledge, asking the LLM to devise a sequence of function calls or actions to reach the goal. The output is typically a list of steps or a structured plan (for example, a list of function names with arguments – effectively an LLM-generated program).
- **Executing Plans:** After a plan is created, `POST /planner/execute` will **run the plan** step by step ⁴⁸ ⁴⁹. The request includes the plan (which could be a list of function calls with parameters), and the Planner service will iterate through each function call, invoke it using the Function Caller, and gather the results. The final response includes the outcome of each step (function outputs) in order ⁵⁰ ⁵¹. This two-phase approach (plan then execute) allows for human verification or modification of the plan before running it, and also mirrors how chain-of-thought planning can be separated from action execution.
- **Corpora Awareness:** `GET /planner/corpora` lists all available corpus IDs the planner knows about ⁵². This suggests the Planner can operate in the context of different knowledge bases (for different users, domains, or sessions) and may select information or tools relevant to the specified corpus when planning.
- **Integrating Reflections:** The Planner service also has endpoints to retrieve reflection history and semantic arc for a corpus (`GET /planner/reflections/{corpus_id}` and `/planner/reflections/{corpus_id}/semantic-arc`) ⁵³ ⁵⁴. This is likely used to provide the LLM with context about what has been learned or what insights have been gathered previously, so it can plan with that in mind. There's also a `POST /planner/reflections/` endpoint to add a new reflection by providing a message (prompt) and letting the system generate and store the reflection ⁵⁵ ⁵⁶. This is similar to the Awareness service's reflection addition, suggesting the Planner can trigger reflective thinking as part of the workflow (for example, asking the LLM to analyze or critique the result after executing a plan step).

In essence, the Planner is the **conductor** that leverages the LLM's capabilities to reason about tasks and decide on using tools. It works closely with the LLM Gateway (for the actual reasoning step) and with the Function Caller (for executing chosen actions). By maintaining awareness of corpora and reflections, it ensures that plans are informed by the AI's knowledge base and past insights.

LLM Gateway

The **LLM Gateway** is a service that acts as a **proxy/facade to the underlying large language model (LLM)**, especially one that supports OpenAI-style function calling ⁵⁷. This gateway allows FountainAI to interface with different LLM providers or models in a consistent way. Its main responsibilities include:

- **Handling Chat Requests:** The primary endpoint is `POST /chat`, which is used to **conduct a chat session or query with the LLM** ⁵⁸. The request typically contains the conversation context (history of messages), the user's latest prompt or objective, and importantly, a list of available functions (tools) that the LLM can choose to call ⁵⁹ ⁶⁰. The LLM Gateway packages

this information and sends it to the actual LLM (for example, an OpenAI GPT-4 API or another model endpoint).

- **Function-Calling Support:** Because the request includes a **functions list** and an instruction to the model that it can invoke functions, the LLM can decide to output a function call as part of its response. The LLM Gateway captures such function call responses and can route them to the Function Caller. In many cases, the Planner service (or whatever client calls /chat) will loop: sending the user query, getting either an answer or a function call, executing the call via Function Caller, and then sending the result back to the LLM via this gateway, until the LLM produces a final answer. The `function_call` field in the request can also be set to "auto" to let the model decide when to use a function ⁶¹.
- **Model Agnosticism:** The gateway is designed to work with any LLM that supports the function-calling interface, not just a single model ⁵⁷. This means FountainAI could be configured to use different models (OpenAI, Azure, local models, etc.) without other services needing to know the specifics. The /chat endpoint ensures the prompt, functions, and parameters are formatted correctly for the model.

In summary, the LLM Gateway **delegates the reasoning** to the AI model given a user's objective and toolset ⁵⁹. It abstracts the details of the LLM API and provides a streamlined interface for the Planner (or other components) to get model-generated decisions and responses. By including tool information in the prompt, it enables the model to perform *augmented reasoning* – i.e., use tools when necessary to fetch information or perform actions.

Persistence Layer

The **Persistence Service** is the backbone that **stores and indexes all of FountainAI's knowledge artifacts** using a Typesense engine ¹. All data related to corpora – baselines, drifts, reflections, functions, history – is persisted via this service so that it can be queried and retrieved efficiently by the other components. Key aspects of the persistence layer include:

- **Corpus Storage:** The service enables creation of new corpora (`POST /corpora`) and listing existing corpora (`GET /corpora`) ⁶² ⁶³. A corpus is identified by a unique ID and acts as a namespace for all associated data (this prevents, for example, tools or reflections from different projects or users from mixing together).
- **Baseline and Drift Storage:** When a baseline snapshot is added to the system, the persistence service stores it under the corresponding corpus (`POST /corpora/{corpusId}/baselines`) ⁶⁴ ⁶⁵. This includes the baseline content and an ID for that baseline version ⁶⁶ ⁶⁷. Drift documents or analyses would similarly be stored, typically as separate entries (often with a naming convention linking them to a baseline). By saving baselines and drifts, the platform builds a time-indexed record of knowledge states, which can later be searched or analyzed.
- **Function Indexing:** All function (tool) definitions registered via the Tools Factory are ultimately saved through the persistence layer as well. The endpoint `POST /corpora/{corpusId}/functions` adds a function record to a corpus ⁶⁸ ⁶⁹. There are also global endpoints to list all functions (`GET /functions`) and get details of a specific function (`GET /functions/{functionId}`) from the persistent store ⁷⁰ ⁷¹. This ensures the Function Caller and Planner can quickly lookup available functions. Each stored function entry includes its unique ID, name, description, and the HTTP method/path for invocation ⁷² ⁷³.
- **Reflection and History Storage:** Reflections generated by the AI are stored via `POST /corpora/{corpusId}/reflections` ⁷⁴ ⁷⁵, which records the reflection question and answer content under the corpus. The persistence service also supports retrieving all reflections with pagination (`GET /corpora/{corpusId}/reflections`) ⁷⁶ ⁷⁷. Additionally, although not shown in the snippet above, the history of changes (baselines added, drifts, reflections, etc.)

for each corpus would be stored and could be retrieved to support the history timeline and semantic arc features.

Because it is **Typesense-backed**, the persistence layer likely provides not just simple storage but also **semantic indexing and search** capabilities. This means the AI can perform semantic lookups – for example, finding which past reflection or baseline is relevant to a new query – using vector search or full-text search. By maintaining a structured yet searchable memory, the persistence service allows FountainAI to scale its knowledge and retrieve relevant information when reasoning.

Platform Architecture and Data Flow

FountainAI's architecture is composed of these modular services that work in concert to achieve intelligent behavior. The typical data and task flow in the platform might look like the following:

1. **Bootstrapping an AI Agent:** When starting a new session or agent, a client (or the system) calls the **Bootstrap Service** to create a new corpus and seed it. The bootstrap process will coordinate with the **Baseline Awareness Service** to initialize the corpus and with the **Roles/Reflection** subsystems to add default GPT roles and a health-check reflection ¹⁹ ²⁰. This sets up a fresh knowledge base and some initial self-monitoring capability for the AI.
2. **Registering Tools:** If there are external functions or APIs the AI might use, they are registered via the **Tools Factory**, which stores their definitions in the **Persistence Layer** (Typesense index). Now the AI's "toolbox" is populated with functions it can call, and the **Function Caller** service is aware of these tools ³ ⁹. (Tools can also be added or updated on the fly as the system is running.)
3. **User Objective and Planning:** The user provides an objective or query to the **Planner Service**. The Planner then prepares a prompt that includes the user's request and contextual information (relevant baseline knowledge, recent history or reflections, and the list of available functions) ⁵⁹. It sends this to the **LLM Gateway** via a `/chat` request, delegating the reasoning to the LLM ⁷⁸. The LLM, seeing the objective and the tool options, may respond with a proposed plan or directly with a function call as the next step (thanks to the function-calling interface). For example, the LLM might decide: "To answer this, I should call the `search_tool` function with argument X".
4. **Function Orchestration:** If the LLM's response includes a function call, the Planner/LLM Gateway will invoke the **Function Caller** to execute that call. The Function Caller looks up the function's details (endpoint, method, etc.) and performs the HTTP request or operation ¹⁰. The result (e.g., the API response) is returned to the LLM via the Gateway. The LLM then continues reasoning with that new information. This loop may repeat, with the LLM calling multiple tools in sequence (this is the *orchestration* process), until it arrives at an answer or a complete plan.
5. **Completing the Plan:** Once the LLM has formulated a full solution or all necessary steps, the Planner service compiles this into a structured **Plan** (a series of steps). If not already executed during the reasoning loop, the client can call `planner/execute` to run any remaining steps through the Function Caller ⁵⁰. The outputs of each step are collected and delivered.
6. **Reflection and Learning:** After or during execution, the system can generate **reflections** to evaluate the outcome. The Planner or Bootstrap might call the Awareness service (or use its own `/planner/reflections` endpoint) to pose questions like "Did the plan succeed? What could be improved?" to the LLM. These reflections (and their answers) get stored in the **Baseline Awareness** (and persistence) as part of the corpus's history ³⁶ ⁷⁴. Over time, a series of reflections provides a form of *automated feedback loop* where the AI learns from each attempt or update.
7. **Updating Knowledge – Baseline & Drift:** Whenever new information is introduced (e.g., the user uploads new data or the world changes), it can be added as a new **baseline** via the

Bootstrap or Awareness service. Upon adding a baseline, the **drift analysis** kicks in to highlight changes compared to previous knowledge ³⁰. The **narrative patterns** analysis might reveal new themes or shifts in tone. All these are stored and the **semantic history** is updated so that future queries to the LLM can leverage the historical context and avoid repeating past mistakes or omissions.

8. **Iterative Improvement:** The next time the user asks a question or sets an objective, the Planner and LLM will have access to a richer context: a full history of what happened, prior solutions, reflections on what worked or failed, and updated knowledge. The AI can reason semantically about this context (using the Awareness service's summaries or semantic arc if needed) to produce better plans. In some cases, a reflection might suggest creating a new tool or a new role – which can then be fed back into the Tools Factory or Bootstrap (role promotion) to extend the system's capabilities dynamically ²⁴.

All components communicate typically over RESTful APIs (as evidenced by their OpenAPI specs). The architecture is microservice-based, but all parts are orchestrated to give the effect of an intelligent, self-improving agent. The **Persistence layer** (with Typesense) ensures that any component can quickly query the stored knowledge (for example, the Planner might search past reflections or baselines to decide which tools to use or what the user might really need). This design allows FountainAI to be **extensible** (new tools or models can be integrated), **context-aware** (through persistent memory of past interactions and content), and **robust** (by iterating plans with reflection and detecting drift to adapt to new information).

Key Use Cases and Features

FountainAI's unique architecture supports several advanced use cases and features that set it apart from a basic LLM interface. Some of the main capabilities include:

- **Semantic Reasoning and Knowledge Summarization:** The platform can perform deep **semantic analytics** on the information in a corpus. By using baselines, drift documents, and narrative pattern analysis, FountainAI can understand and summarize how a body of knowledge changes over time. For example, it can provide a **semantic summary of the corpus history** or a "*semantic arc*" that highlights the evolution of themes and insights ⁷⁹ ⁴³. This helps in reasoning about the context; the AI can explain what has been happening or how a situation developed, which is crucial for domains like personal coaching, long-term projects, or storytelling.
- **Dynamic Function Orchestration:** FountainAI allows LLMs to **orchestrate external functions** as part of their reasoning. Through the Tools Factory and Function Caller, an LLM can access a library of operations (APIs, database queries, calculations, etc.) and invoke them in real time to get information or perform actions ¹⁶. This means the AI is not limited to its trained knowledge; it can, for instance, call a web search API, fetch user data, or execute computations on the fly. The Planner service specifically uses this to break objectives into **function call plans**, enabling complex workflows to be automated by the AI (this is akin to having the AI write and execute code to solve a problem).
- **Automated Reflection and Self-Improvement:** A standout feature of FountainAI is its ability to perform **automated reflection**. After completing tasks or at certain intervals, the AI can generate reflections – effectively asking itself questions like "What did I learn?", "What went wrong?", or "How can I do better next time?" – and then answering them using the LLM. These reflections are stored in the corpus ³⁶ ⁷⁴. They serve two purposes: (1) to provide transparency and explanations of the AI's thought process (useful for users to understand the AI's reasoning), and (2) to create a feedback loop where the AI's future planning can avoid past errors. Moreover,

FountainAI can even turn insightful reflections into new **role prompts** (using the Bootstrap role promotion) – effectively **learning new personas or approaches** to improve performance ²⁴ .

- **Drift Analysis and Change Detection:** The platform is designed to handle scenarios where information updates or evolves. **Drift analysis** automatically detects and records how a new piece of information (a new baseline) differs from previous knowledge ²⁸ . This is crucial in long-running applications (like coaching or monitoring systems) where the situation may change gradually – the AI remains aware of what’s new or what trend is forming. By logging drifts and patterns, FountainAI can alert users to important changes or adapt its advice and plans accordingly. For example, in a business coaching context, if new sales data shows a shift (drift) in customer behavior, the AI would note that and adjust its recommendations.
- **Structured Multi-Step Planning:** Instead of answering questions in one go, FountainAI emphasizes **step-by-step planning** for complex objectives. The LLM (via the Planner) can output intermediate steps or sub-goals, possibly with function calls at each step, to systematically approach a problem ⁸⁰ ⁵⁰ . This capability is vital for complex tasks (e.g., “*Help me analyze my financial data and draft a report*”) where the solution involves multiple actions (gather data, analyze trends, generate text) – the AI can plan these actions, execute them, and adjust as needed. It’s essentially the AI doing problem decomposition and execution autonomously.
- **Contextual Adaptation with Roles:** With its concept of **GPT roles**, FountainAI can inject different perspectives or skills into the AI’s responses. The default roles seeded during bootstrap give the AI a balanced starting persona set, but as the system encounters new challenges, it can create or learn new roles. This could be seen as having multiple experts in one AI – e.g., a “Planner” role, a “Critic” role, a “DomainExpert” role – and switching between them or consulting them internally. This role-based reasoning makes the AI’s output more robust, as it can self-check or iterate through various approaches before presenting a final answer.

Overall, FountainAI is built for **robust, explainable, and adaptive AI interactions**. It goes beyond single-turn Q&A by maintaining a memory of past interactions and knowledge, by enabling tool use for extended capabilities, and by incorporating mechanisms for self-reflection and adjustment. This makes it suitable for applications like personal coaching agents, complex decision-support systems, research assistants, or any scenario where an AI needs to **continuously learn and reason** in a changing environment. The combination of its core components ensures that the AI can plan actions, execute them, learn from the outcomes, and refine its knowledge, aligning with the platform’s mission of delivering a thoughtful and evolving AI assistant ¹⁶ ³¹ .

1 62 63 64 65 66 67 68 69 70 71 72 73 74 75 76 77 persist.yml

file://file-TRqbobR7N5nnLYBFVtBrrj

2 3 4 5 6 7 8 tools-factory.yml

file://file-LM2mLuM9u824cSCq6a7zc5

9 10 11 12 13 14 15 16 function-caller.yml

file://file-FdchupWBgtkS5mdyD1kmdL

17 18 19 20 21 22 23 24 25 26 27 28 29 30 bootstrap.yml

file://file-DmYhEPkgHVt6irpdeuDKFN

31 32 33 34 35 36 37 38 39 40 41 42 43 44 79 baseline-awareness.yml

file://file-U4Nr8WUFJLz9Y5LfZTMPvo

45 46 47 48 49 50 51 52 53 54 55 56 80 planner.yml

file://file-KjtBPERcsKwyoZZyUPQNqx

57 58 59 60 61 78 llm-gateway.yml

file://file-42vkSwoNDubFckjxsVFAYV